

Model fitting

Katie Schuler

2023-10-05

 Under Construction

0.1 You are here

0.2 Model building

0.3 Linear model review

- how would we specify this model
- how would we write it in R
- estimate the free parameters

0.4 Estimate free parameters

```
# with base R  
lm(y ~ 1 + x, data = data)
```

Call:

```
lm(formula = y ~ 1 + x, data = data)
```

Coefficients:

(Intercept)	x
0.6	0.7


```
# with infer workflow
data %>%
  specify(y ~ 1 + x) %>%
  fit()
```

```
# A tibble: 2 x 2
  term      estimate
  <chr>      <dbl>
1 intercept 0.600
2 x         0.7
```

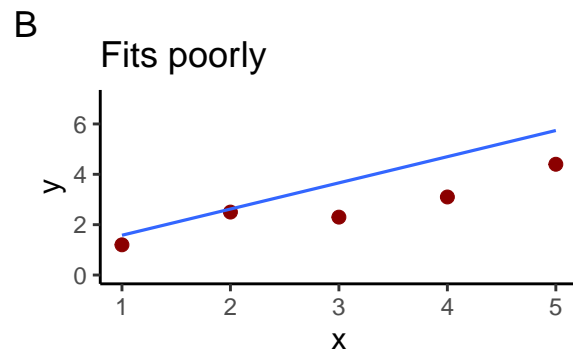
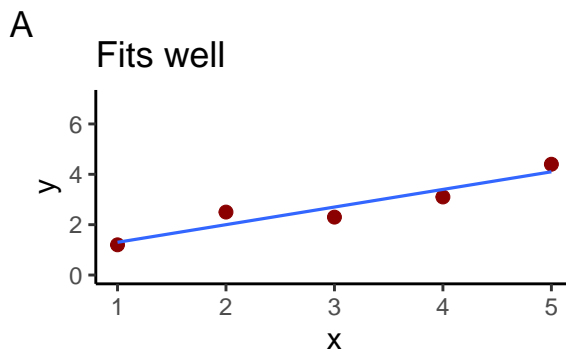
0.5 Model fitting basics

Linear model

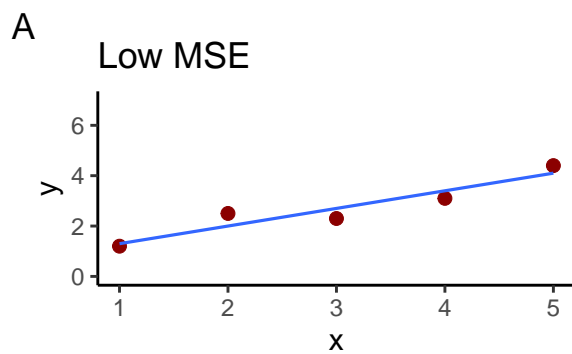
```
# A tibble: 5 x 2
      x     y
  <dbl> <dbl>
1     1  1.2
2     2  2.5
3     3  2.3
4     4  3.1
5     5  4.4
```



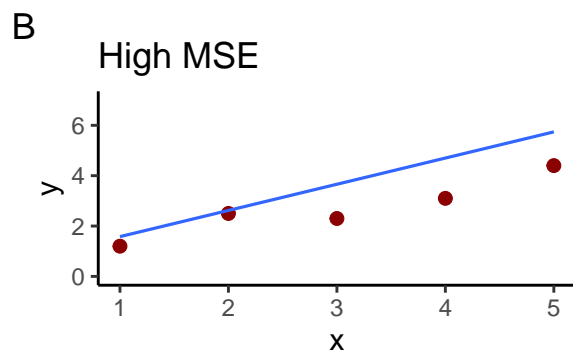
We can see which fits better with our eyes.



```
mean(mseA$sq_err)
```



x	y	pred	err	sq_err
1	1.2	1.3	-0.1	0.01
2	2.5	2.0	0.5	0.25
3	2.3	2.7	-0.4	0.16
4	3.1	3.4	-0.3	0.09
5	4.4	4.1	0.3	0.09

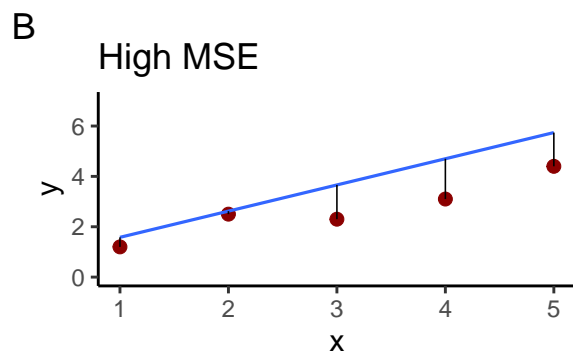
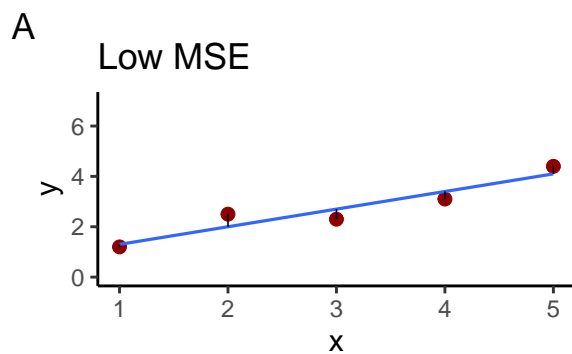


x	y	pred	err	sq_err
1	1.2	1.58	-0.38	0.1444
2	2.5	2.62	-0.12	0.0144
3	2.3	3.66	-1.36	1.8496
4	3.1	4.70	-1.60	2.5600
5	4.4	5.74	-1.34	1.7956

```
[1] 0.12
```

```
mean(mseB$sq_err)
```

```
[1] 1.2728
```



```
lm(y ~ 1 + x, data = data)
```

Call:

```
lm(formula = y ~ 1 + x, data = data)
```

Coefficients:

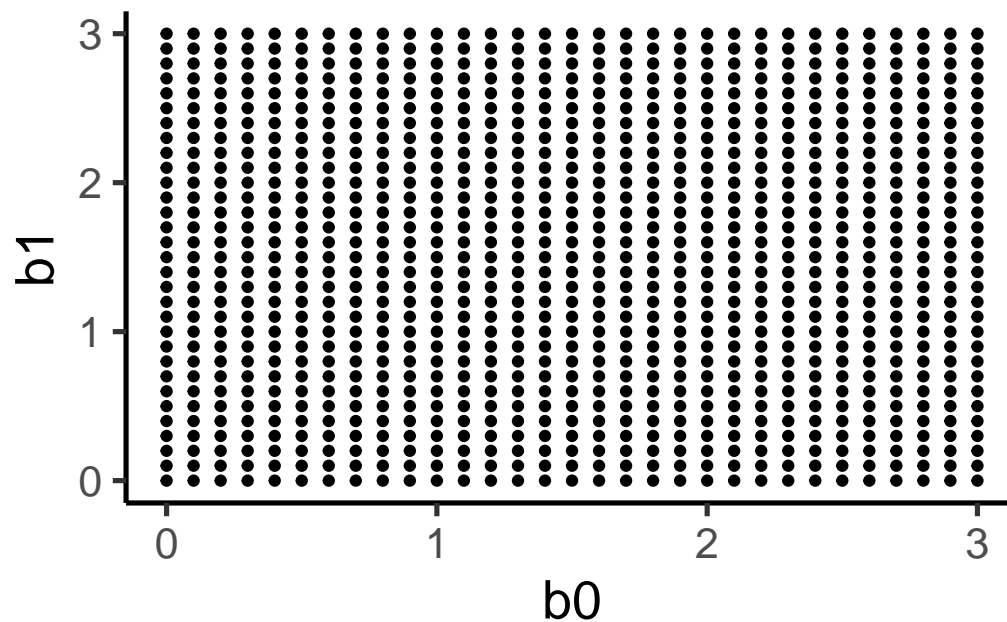
(Intercept)	x
0.6	0.7

```
data %>%  
  specify(y ~ 1 + x) %>%  
  fit()
```

```
# A tibble: 2 x 2  
  term      estimate  
  <chr>      <dbl>  
1 intercept 0.600  
2 x         0.7
```

```
b0 <- seq(from = 0, to = 3, by = 0.1)  
b1 <- seq(from = 0, to = 3, by = 0.1)  
possible_weights <- expand.grid(b0 = b0, b1 = b1)
```

```
ggplot(data = possible_weights,  
  mapping = aes(x = b0, y = b1)) +  
  geom_point()
```



```

# compute the sum of squares for those weights on a dataframe
sum_squares <- function(b0, b1) {

  data %>%
    mutate(pred = b0 + b1*x) %>%
    mutate(err = pred-y) %>%
    mutate(sq_err = err^2) %>%
    select(sq_err) %>%
    sum()

}

error_surf <- possible_weights %>%
  rowwise() %>%
  mutate(sum_sq = sum_squares(b0, b1)) %>%
  ungroup

error_surf

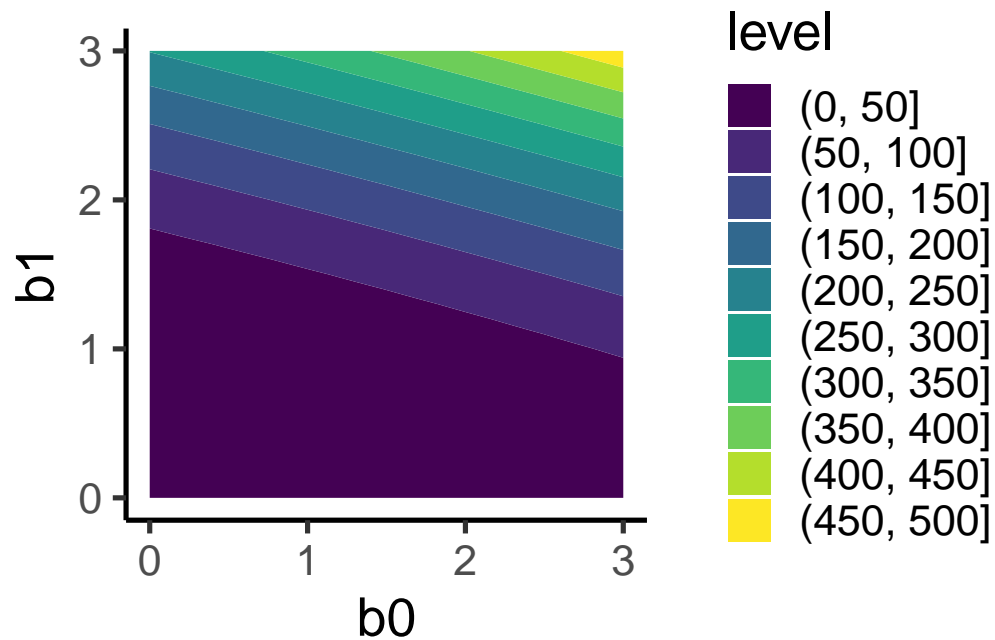
# A tibble: 961 x 3
   b0    b1 sum_sq
<dbl> <dbl> <dbl>
1     0     0  42.0
2    0.1     0  39.3
3    0.2     0  36.8
4    0.3     0  34.3
5    0.4     0  32.0
6    0.5     0  29.7
7    0.6     0  27.6
8    0.7     0  25.5
9    0.8     0  23.6
10   0.9     0  21.7
# i 951 more rows

error_surf %>% filter(sum_sq < 0.608)

# A tibble: 1 x 3
   b0    b1 sum_sq
<dbl> <dbl> <dbl>
1  0.6  0.7  0.600

ggplot(error_surf, aes(b0, b1, z = sum_sq)) +
  geom_contour_filled()

```



```
# %>%
# sum(.$sq_err)
# summarise(sum_sq = sum(sq_err)) %>%
```

```
# return sum of squares as a column next to
```

```
# mse <- function(data, b0, b1) {
#   model_value <- b0 + b1*data[1]
#   resid <- data[2] - model_value
#   sq_err <- resid^2
#   sum(sq_err)
# }
```

```
# possible_weights %>% mutate(
#   mse = mse(1, 1, b0, b1)
# )
```

- in context of model building more broadly
- a general overview of the concept

0.6 Mean squared error

Cost function.

0.7 Error surface

- We can visualize the error surface for simple example: 2 parameters, β_0 and β_1 , and the cost function (mean square error).
- Show nonlinear model v linear model figs
- goal is to find the minimum point
- notice the nonlinear model can have local minimums but lm has only 1. Because lm is a **convex** function.

0.8 Gradient descent

IF we want to estimate the free parameters in a way that would work broadly, for linear or nonlinear models, we can use **gradient descent**.

- machine learning / optimization.
- If we have a lot of data, we could use **stochastic gradient descent** which is the same except we...

0.9 Ordinary least squares

As we saw above, linear models have the special property that they have a solution, the OLS. Rather than searching the error surface iteratively via gradient descent (optimization), we can solve for this point directly with **linear algebra**.

- matrix approach: we write the 3-step function.
- use `lm()` in R.
- infer approach:
 - `specify()`, `fit()`

0.9.1 Further reading

- [Ch. 8 Fitting models to data](#) in Statistical Modeling